

TPDP, WPES and CCS talks

Sebastian Meiser¹

1: University College London, United Kingdom, e-mail: s.meiser@ucl.ac.uk

October 19, 2018

Abstract

This is a document of summaries of talks I attended at CCS 2018. Please note that in most cases I have not read the paper before writing the summary, so my summaries might be very shallow and sometimes factually incorrect. If you notice any inconsistencies, errors or would like me to include aspects that I've missed, please just send me an email.

Contents

1	Pre-Conference Workshops	3
1.1	Invited talk TPDP: Composition, Verification, and Differential Privacy	3
1.1.1	Formally verifying privacy	3
1.1.2	Advanced composition	3
1.2	Modularity	3
1.3	Invited Talk TPDP 2: Deploying Differential Privacy for Learning on Sensitive Data	3
1.3.1	Private neural networks	4
1.4	Local differential privacy for evolving data	4
1.5	WPES: TightRope: Towards Optimal Load-balancing of Paths in Anonymous Networks	4
1.6	WPES: ClaimChain: Improving the Security and Privacy of In-band Key Distribution for Messaging	5
1.7	WPES: What's a little leakage between friends? (Short)	5
1.8	WPES: DynaFlow: An Efficient Website Fingerprinting Defense Based on Dynamically-Adjusting Flows (Short)	6
2	CCS, Tuesday, Privacy	6
2.1	ABY3: A Mixed Protocol Framework for Machine Learning	6
2.2	Voting: you can't have privacy without verifiability	7
3	CCS, Tuesday, Differential Privacy 2	7
3.1	Preserving Both Privacy and Utility in Network Trace Anonymization	7
3.2	Toward Detecting Violations of Differential Privacy	8
3.3	Secure Computation with Differentially Private Access Patterns	8
3.4	DP-Finder: Finding Differential Privacy Violations by Sampling and Optimization	8
4	CCS, Wednesday, Cyberphysical Systems	9
4.1	Scission: Signal Characteristic-based sender identification and intrusion detection in automotive networks	9
4.2	Detecting Attacks Against Robotic Vehicles: A Control Invariant Approach	9
4.3	Truth Will Out: Departure-Based Process-Level Detection of Stealthy Attacks on Control Systems	10
4.4	On the Safety of IoT Device Physical Interaction Control	10

5	CCS, Wednesday, Usable Security	11
5.1	Asking for a Friend: Evaluating Response Biases in Security User Studies	11
5.2	Towards Usable Checksums: Automating the Integrity Verification of Web Downloads for the Masses	12
5.3	Investigating System Operators' Perspective on Security Misconfigurations	12
5.4	Detecting User Experience Issues of the Tor Browser In The Wild	13
6	CCS Thursday, Web Security 1	14
6.1	Predicting Impending Exposure to Malicious Content from User Behavior	14
6.2	Clock Around the Clock: Time-Based Device Fingerprinting	14
6.3	Web's Sixth Sense: A Study of Scripts Accessing Smartphone Sensors	15
7	CCS Thursday, Web Security 2	16
7.1	Mystique: Uncovering Information Leakage from Browser Extensions	16
7.2	How You Get Bullets in Your Back: A Systematical Study about Cryptojacking in Real-world	16
7.3	MineSweeper: An In-depth Look into Drive-by Cryptocurrency Mining and Its Defense . . .	17
7.4	Pride and Prejudice in Progressive Web Apps: Abusing Native App-like Features in Web Applications	18
8	CCS, Thursday, Tor	18
8.1	Deep Fingerprinting: Undermining Website Fingerprinting Defenses with Deep Learning . . .	18
8.2	Privacy-preserving Dynamic Learning of Tor Network Traffic	19
8.3	DeepCorr: Strong Flow Correlation Attacks on Tor Using Deep Learning	20
8.4	Measuring Information Leakage in Website Fingerprinting Attacks and Defenses	21
9	Post-conference Workshop AISec	21
9.1	Hate speech detection	21
9.2	Towards Query Efficient Black-box Attacks: An Input-free Perspective	22
9.3	Stochastic Substitute Training: A General Approach to Craft Adversarial Examples against Defenses which Obfuscate Gradients	22
9.4	Adaptive Grey-Box Fuzz-Testing with Thompson Sampling	23
9.5	Toward Smarter Vulnerability Discovery Using Machine Learning	23

1 Pre-Conference Workshops

1.1 Invited talk TPDP: Composition, Verification, and Differential Privacy

Speaker: Justin Hsu

After introducing differential privacy in general, we see a few well-known composition results. Post-processing is seen as an instantiation of sequential composition of an (ϵ, δ) -DP mechanism with a $(0, 0)$ -DP mechanism, which is kind of nice. Similarly, local DP is presented as an instantiation of parallel composition.

1.1.1 Formally verifying privacy

Our goals here are twofold: we want to have *dynamic* verification, i.e., raise errors if DP is violated and *static* verification, i.e., checking DP on all possible inputs. As a side-note, we want to simplify verification by using composition results; composition allows verification techniques to have much better automation.

The Fuzz family of languages allows for describing each type with a metric and, using group privacy ($(\epsilon, 0)$ -DP for $\Delta_f = 1$ implies $(k \cdot \epsilon, 0)$ -DP for $\Delta_f = k$). The description allows for an immediate static analysis that can use both sequential and parallel composition. These languages, so far, cannot deal well with ADP, but the speaker is currently working on that.

Privacy as an approximate coupling Idea: verify privacy by game hops. We somehow relate pairs of sampling instructions to show properties of a pair of inputs, but I'm not sure how exactly we do that. This static-analysis technique seems to be related to how Tim and me modeled differential indistinguishability, i.e., as a property of a pair of programs/machines. They, so far, cannot automatically generate proofs, but can check them.

1.1.2 Advanced composition

They analyze the old “advanced composition theorem” by Dwork, Rothblum and Vadhan. The speaker says that the analysis is more complicated, since the composition theorem needs to be applied “as a block”, instead of on small units of the program. The speaker then mentions that novel approaches (e.g., by Mironov on RDP) makes formal verification easier, as it allows to analyze a program step-by-step. I think that using our privacy buckets approach, such an analysis could be improved even more and made flexible as well.

1.2 Modularity

The verification approach aims at being able to modularly compose differential privacy mechanisms. Their approach does not solve a fundamental problem that black-box composition of DP mechanisms have: after black-box composition there is potentially too many points where noise is added. So, a more thorough approach would divide DP mechanisms into summarizing data (e.g., via statistical queries) and adding noise. Then, composition would compose the summarization operations and then only add noise once. Such an approach would deteriorate utility much less.

1.3 Invited Talk TPDP 2: Deploying Differential Privacy for Learning on Sensitive Data

Speaker: Ulfar Erlingsson

The speaker mentions problems with RAPPOR, mostly that for the quantities of data they require, the privacy guarantees are deteriorating too fast. The main problem seems to be that they have to use local DP mechanisms.

The aim of their Prochlo realization is to combine the local differential privacy mechanisms with some central DP approach.

They notice that if they don't need correlations between data points from the same user, anonymous communication can help them. As a result, they have a partially central approach. They group data based on useful features.

Their approach combines three aspects of uncertainty: the uncertainty introduced by the local DP mechanism (via randomized response), the uncertainty introduced by the combination (anonymity and random shuffling) and the uncertainty applied to the results afterwards (in a centralized DP manner).

1.3.1 Private neural networks

The speaker identifies the problem of machine learning models, e.g., NNs, memoize the training data. This confirms known results on adversarial ML and works by Dawn Song and Vitaly Shmatikov.

Specifically, they introduce canaries, i.e., specific numbers, and want to check whether the NN is surprised to see this canary or not. In other words, they check whether the NN learnt this canary. Their experiments show that NNs indeed learn such canaries. Learning the canaries seems to reach its worst case early, with the convergence of the model.

The speaker emphasizes that this memoization is different from overfitting, because in spite of memoization generalization works well.

The speaker emphasizes that there are edge cases that the NNs seemingly have to memorize exceptions, because there does not seem to be a rule to characterize these edge cases. These exceptions are not necessarily important for the final model, as they contradict the general paradigm of machine learning: to learn general properties of samples and not exceptions. This makes differential privacy a natural fit for ML.

We now get an introduction into the privacy-preserving stochastic gradient descent (SGD) [?] work and the PATE paper [?].

In their case study, they use prior knowledge and use a strong bayesian prior.

1.4 Local differential privacy for evolving data

Speaker: Matthew Joseph, collaboration with Aaron Roth, Jonathan Ullman and Bo Waggoner

Paper: <https://arxiv.org/pdf/1802.07128.pdf>

The main difficulty they tackle is evolving data, i.e., data that can change over time. The simplest solution in terms of differential privacy would be to simply have a randomized response for every user in every round. While this is very simple, privacy deteriorates over time. To moderately improve it, you can apply subsampling, i.e., ask different subsets of users every time. This is better in terms of privacy, but the error is still large. Local sparse vectors (?) would allow for nice privacy guarantees, but require an exponentially larger number of samples.

The solution lets users “vote” to change statistics, i.e., they construct an adaptive mechanism.

In each epoch t , each user creates a localized bit, then generates a vote for or against updating and then depending on the average vote the analyst decides whether or not the estimate should be updated (requiring more interaction) or not.

Each user can guess whether or not the global estimate has changed significantly. We aren’t completely sure on the maths here, but apparently each user locally simulates what the estimate should be (in the current round) and compare that to the previous global estimate. If these don’t differ much, the user is content; if they do differ significantly, the user votes for a change.

If a change occurred, the analyst collects the changed estimates from the users to update the global estimate. Otherwise the analyst just keeps the previous estimate.

I’m not completely convinced yet by the proof intuition for why their mechanism satisfies differential privacy; a closer look into their paper ¹ might be in order to understand what exactly they are doing and why that is privacy-preserving, particularly towards the data analyst that collects the votes and combines the estimates.

1.5 WPES: TightRope: Towards Optimal Load-balancing of Paths in Anonymous Networks

Speaker: Hussein Darir (University of Illinois at Urbana-Champaign), collaboration with Hussein Sibai (University of Illinois at Urbana-Champaign), Nikita Borisov (University of Illinois at Urbana-Champaign), Geir Dullerud (University of Illinois at Urbana-Champaign), and Sayan Mitra (University of

¹available online: <https://arxiv.org/pdf/1802.07128.pdf>

Illinois at Urbana-Champaign)

The paper is about Tor; some Tor circuits tend to be slow. Knowing the current state of the network (in terms of congestion and usage) leaks too much information, but maybe differential privacy can help us here. The goal is to use load balancing with locally optimal mechanisms and this mechanism should then be differentially private.

They assume that each user only holds a single static circuit that is active all the time.

Each relay computes the ratio between their capacity over the number of paths going through it. Then they choose the relay with the minimal ratio, called the bottleneck relay (it is most congested). The path(s) going through the bottleneck relay are tagged with the ratio of the bottleneck relay. Then from the other relays, the capacity used by the path through the bottleneck relay is subtracted from these relays' capacities, and then this path is removed. The algorithm then is repeated to compute the bandwidth available to each path. They evaluate this and find (unsurprisingly) that the bandwidths are unfairly distributed over paths.

Their (non-private) algorithm now takes the existing bandwidth capacities and ratios into account and also computes the ratio where one more path is added to the relay. They find the most congested node, i.e., the one where the ratio would be worst if we added one. We then remove this relay (and paths through it; and update the ratios). This process is repeated until we can build a path. Using this technique makes path selection much more fair, but violates privacy.

To achieve differential privacy, they need up-to-date statistics. I'm not completely sure how they achieve that; they seem to sample many times, i.e., have each user locally simulate how the network might have changed? I'm not quite sure what exactly they do and what privacy guarantees they achieve, but their evaluation shows that still some improvement over a simple random choice (without optimization) is possible.

They did not analyze the impact their work has on the anonymity provided by Tor; just the impact of releasing their histogram.

1.6 WPES: ClaimChain: Improving the Security and Privacy of In-band Key Distribution for Messaging

Speaker: Wouter Lueks (Ecole Polytechnique Fédérale de Lausanne), collaboration with Bogdan Kulynych (Ecole Polytechnique Fédérale de Lausanne), Marios Isaakidis (University College London), George Danezis (University College London), and Carmela Troncoso (Ecole Polytechnique Fédérale de Lausanne)

Paper: <https://arxiv.org/abs/1707.06279>

The classical PKI problem is: how does Alice find Bob's key? They analyze how to distribute keys for an email structure that allows for encrypted communication. Whenever someone sends an email, they want to attach some small datastructure that includes claims about their own key and about their friends' keys. Moreover, they want to hide whose keys they are communicating from unauthorized people. They intend to do that by adding some access control mechanism, specifying who can access these records. Finally, they want to get non-equivocation, i.e., sending around different keys supposedly belonging to the same entity, accessible by different people.

Their Approach They store all claims (of the form "entity, key") into an encrypted dictionary. To hide the indexes of the dictionary, they put a verifiable random function in there, i.e., some apparently random value; they communicate this value to people of interest. They then put a proof that they computed the VRF correctly into the encrypted part. They then chain these claims together with a hash chain. This basically is the ClaimChain that can be attached to emails.

1.7 WPES: What's a little leakage between friends? (Short)

Speaker: Sebastian Angel, collaboration with David Lazar, Ioanna Tzialla

Paper: <https://arxiv.org/abs/1809.00111>

Metadata-private messaging systems allow communication without leaking metadata to providers, servers or users not involved in the communication. In the simplest case, everyone sends packets to everyone else. Existing MPM systems avoid this broadcast by limiting the number of messages per round. Clients now

have to start scheduling their own communication (and potentially wait until their messages are being sent). To actually communicate, you can have a dialing round in which you negotiate a communication round.

They then analyze what happens if a malicious user tries to gain information about whether others are talking: the idea is that an adversary can ask a user whether they want to talk; since every user only has a limited number of (real) messages per round, they can see whether the user is “free” or “already in conversation”. To counter this, they envision a private answering machine, that hides from callers whether or not a user is talking; moreover eventually a caller needs to get through and it should be efficient, i.e., the capacity of the answering machine should be significantly smaller than full broadcast. Their proposal has several limitations, including that everyone needs to have a max number of friends m ; then statically map callers to rounds (mod m). Drawbacks are potential limitation and leakage of the number of friends and an increased latency (if the answering machine has low capacity or the user many friends).

1.8 WPES: DynaFlow: An Efficient Website Fingerprinting Defense Based on Dynamically-Adjusting Flows (Short)

Speaker: David Lu (MIT PRIMES), collaboration with Sanjit Bhat (MIT PRIMES), Albert Kwon (MIT), and Srinivas Devadas (MIT)

Paper: <https://dl.acm.org/citation.cfm?id=3268960>

The paper aims at countering website fingerprinting, particularly when people are using Tor. As we know, Tor is vulnerable to traffic analysis attacks, including website fingerprinting attacks. The talk considers an open-world scenario for fingerprinting. Existing defenses, e.g., supersequence defenses over-approximate all websites within an anonymity set and make all these websites look the same. Similarly, constant flow defenses enforce a constant transmission rate, which introduces a significant overhead as well.

Their approach morphs traffic into fixed bursts: o outgoing and i incoming packets, for fixed values o and i . Moreover, they vary the inter-packet timing interval: t changes every b bursts, t is chosen from some set $\{t_1, \dots, t_k\}$ and up to a adjustments are allowed.²

They evaluate their approach against a couple of existing attacks, using both a “medium security” and a “high security” setting. Their results look promising: the overheads are only 28% of time overhead and about 110% of bandwidth overhead.

2 CCS, Tuesday, Privacy

2.1 ABY3: A Mixed Protocol Framework for Machine Learning

Authors: Payman Mohassel (Visa), Peter Rindal (Oregon State University)

Paper: <https://eprint.iacr.org/2018/403.pdf>

Three party SMPC; each party encrypts their data and sends some of their shares to the other parties. They use three types of sharing: arithmetic (sums; $x = x_1 + x_2 + x_3$), binary (xor $x = x_1 \oplus x_2 \oplus x_3$) and Yao’s garbled circuits.

Their main focus is on machine learning, for which they represent floating point numbers as a pair of integers that they process separately, which is pretty straight-forward. They use a trick in which they briefly fall back to a 2-out-of-2 secret sharing mechanism for performing multiplications securely without adding a lot of overhead. Moreover, they can deal with matrix multiplication with less communication overhead than state-of-the-art.

They show a general way to convert freely between arithmetic secret sharing, boolean secret sharing, and Yao’s garbled circuits.

Finally, they can perform linear and logistic regression as an SMPC and repeat the process to train neural networks. Particularly for neural networks, their performance is orders of magnitude better than previous work on two-party SMPC (called “SecureML”).

²I’m not sure within which time frame a is measured

2.2 Voting: you can't have privacy without verifiability

Speaker: Joseph Lallemand (CNRS, Inria, Loria, Université de Lorraine, France), collaboration with Véronique Cortier (CNRS, Inria, Loria, Université de Lorraine, France)

Paper: <https://hal.inria.fr/hal-01858034/document>

As the title suggests, the paper is about the formal verification of voting. The speaker emphasizes that, particularly in eVoting, there are many potential adversarial parties and attack angles, including dishonest voters, ballot boxes, tallying authorities and communication channels. The main goals they have are privacy (of votes, as an indistinguishability property), verifiability (voters can check that their votes have been correctly cast, including: individuals checking that the vote is in the box, everyone checking that the results have been computed based on the votes in the box and finally, the verification that only valid voters participated).

Privacy and verifiability naturally is contradictory: it is quite easy to achieve one without the other, but hard to achieve both together.

They show, in fact, that their definition of privacy directly implies their definition of individual verifiability. I presume that they implicitly require correctness, as otherwise this is clearly not true.³ They define individual verifiability as the inability of the attacker to modify the result of the voting: The adversary is allowed to cast votes, but must not be able to “remove the honest votes from the result”. I’m not sure what exactly that means, but I think that this is where their implicit correctness assumption comes into play.

They then prove that an attacker that can manipulate individual verifiability (i.e., remove people’s votes) can break privacy. The proof starts with an attacker that can manipulate people’s votes and they then use this attacker to break privacy (just fix Alice’s vote as either 0 or 1) and then, since the **distributions** the adversary has to provide for privacy have to **return the same result** for the voting process, the adversary can learn Alice’s vote by checking whether the tally changed. This proof technique is tailored to the attacker, but in their paper they generalize the technique to work with any attacker.

Learning from their insight, they then define a novel privacy definition based on the verification steps of the protocol. The ballot-box is now dishonest. The attacker is given access to an oracle that lets people verify their votes. The attacker can only see the results after all voters have verified their votes. Their implication still holds.

3 CCS, Tuesday, Differential Privacy 2

3.1 Preserving Both Privacy and Utility in Network Trace Anonymization

Speaker: Meisam Mohammady, collaboration with Lingyu Wang (Concordia institute for information systems engineering), Yuan Hong (Illinois Institute of Technology), Habib Louafi (Ericsson Research Security), Makan Pourzandi (Ericsson Research Security), Mourad Debbabi (Concordia institute for information systems engineering)

Outsourcing network traces is relevant for getting top security monitoring and analytics. Sending network traces to some outsider obviously comes with privacy concerns. The author mentions an anonymization function: the existing prefix preserving anonymization function preserves prefix equality, but has a variety of vulnerabilities. One of these vulnerabilities includes simply injecting a few traces to then find traces for similar subnets.

The adversary is honest-but-curious and tries to find all possible matches between the anonymized and the original traces. Moreover the adversary has α -knowledge, i.e., can have successfully injected α traces.

They send several traces to the analyst and somehow hide the real one in there..? They first hide the original trace, then partition the ptrace, then encrypt each partition a different number of times (thus each partition is prefix-preserved, but globally it isn’t). They create several views then, one of which is the real one (why does that not have the same problem again?). To protect against such madness, they require their fake views to be within a e^ϵ privacy loss of the real view, i.e., they all could have been “real”. I’m not quite

³Consider the protocol that always outputs the same “results” independent of the votes; this preserves privacy, but since the votes aren’t used, they cannot be verified.

sure how they achieve that, but it involves a lot of encrypting and “reverse encrypting”, which might be decryption.

3.2 Toward Detecting Violations of Differential Privacy

Speaker: Yuxin Wang (Pennsylvania State University), collaboration with Ding Ding (Pennsylvania State University), Guanhong Wang (Pennsylvania State University), Danfeng Zhang (Pennsylvania State University), Daniel Kifer (Pennsylvania State University)

The aim of the talk is to test the design of algorithms to find out whether they are differentially private. To this end, they present a tool. They use pure DP in their work. Their aim is to find a good counterexample, i.e., a pair of adjacent databases and a set S s.t. the DP formula is violated. That sounds quite straightforward, but I’m not sure how easy it is to find these counterexamples.

- First they need to find candidate databases; to this end they try to find the most extreme databases that still satisfies DP, such as pairs like $[1, 1, 1, 1]$ and $[2, 2, 0, 0]$ (same sum) or $[1, 1, 1, 1]$ and $[2, 0, 0, 0]$. This kind of strategy is not sound (that’s not their aim anyway), but as long as the tool is meant as a help for programmers, it’s probably a good start.
- For finding a set S given D_0 and D_1 they run the mechanism many times and try to figure out the privacy loss, then selecting the highest ones.
- Given D_0, D_1 and S , they apply a hypothesis test to figure out whether DP is violated: DP being preserved is their null-hypothesis and they try to get a small p-value. I think that makes sense; they have the additional problem that they would need to sample/run the mechanism many times and don’t have a ground truth. Instead they use a clever Fisher-test to get a p-value for whether this particular hypothesis is true.

I like their approach, as it is a nice guide for program designers. The tool runs within 23 seconds, which means it can be used quickly in the development process.

3.3 Secure Computation with Differentially Private Access Patterns

Speaker: Sahar Mazloom (George Mason University), collaboration with S. Dov Gordon (George Mason University)

They look at secure computation; they are looking at secret sharing in a two-party setup. each user performs secret sharing with two servers, but these servers then compute the exact sums in the data. If done naively, we still leak access patterns, in addition to the noiseless results we output. Since generic solutions for hiding access patterns are expensive, they allow a little bit of leakage for these access patterns (i.e., differential privacy). This is related to what Raphael Toledo was working on with his DP Oram project.⁴

The relaxation allows to reduce the asymptotic complexity from $O(n \log n)$ to αn , where α depends on the DP parameters. What they do is that instead of an expensive shuffle operation, they add a number of dummy elements and then perform an oblivious shuffle to mix real and fake elements; each element is annotated with whether it is real or fake (internally, of course).

They can compute a variety of meaningful metrics with their approach including histograms, PageRank and matrix factorization.

3.4 DP-Finder: Finding Differential Privacy Violations by Sampling and Optimization

Speaker: Benjamin Bichsel (ETH Zürich), collaboration with Timon Gehr (ETH Zürich), Dana Drachler Cohen (ETH Zürich), Petar Tsankov (ETH Zürich), Martin Vechev (ETH Zürich)

⁴Mix-ORAM: Using Delegated Shuffles and

The second system for finding differential privacy violations presented at CCS'18. Their approach introduces a lower-bound for pure differential privacy. They notice that what they actually want to do is find some lower bound on the privacy loss (for the whole test S): $\mathcal{L}_{M(D_0)||M(D_1)}^S = \log \frac{M(D_0) \in S}{M(D_1) \in S}$. Once you have found some worst-case inputs and a grasp of the mechanism, this is quite straight-forward; so the main difficulty is finding (approximately) worst-case inputs and then finding out the output distributions. This is important because their work is mechanism-oblivious.

What's more interesting is that they approximate their privacy losses by sampling, which of course makes sense. They then make their privacy losses differentiable (but they don't seem to order them).⁵ They approximate the privacy loss by running $M(D_0)$ many times and $M(D_1)$ many times to then get an approximate for the two probabilities, to then compute the ratio. They do need quite a lot of samples, but I think their sampling approach is pretty interesting.

I think this approach is mostly useful if we're interested in implementation mistakes (or mechanisms that are difficult to analyze mathematically), not if we care about mechanisms for which we know the probabilities (because then we could look at the real privacy loss).

4 CCS, Wednesday, Cyberphysical Systems

4.1 Scission: Signal Characteristic-based sender identification and intrusion detection in automotive networks

Speaker: Marcel Kneib, collaboration with Christopher Huth

The talk is about how physical characteristics can be used for identification of engine control units (ECU) in cars.

Motivation Cars can be attacked via hacks, which can be particularly dangerous when human lives are put at risk (failing brakes, hijacked steering, etc.). The speaker presents the issue intuitively in a case where there is no clear separation between critical components and components with connection to the outside world (internet, bluetooth, etc.). An example for that is the controller area network used for vehicular communication. This system doesn't even have sender authentication for the signals that are sent by components. Apparently it is infeasible to introduce proper integrity mechanisms.

Idea and approach Interestingly, the analog CAN signal (how fast does it rise, how stable is it, etc.) can be used to identify ECU's. Scission now samples the signal, slices it, puts them into three possible groups, extracts features and then performs classification.

Their features include the moments of the function, e.g., mean, standard deviation and variance, skewness, etc.; they require their initial training to be in a safe environment (to avoid poisoning attacks). They also share keys between the ECUs and their module. Intrusion detection now means running the classifier to identify the sender; each identifier is only used by one ECU and in case of a conflict, an alarm is raised. To reduce false-positives, they include a confidence bar and only raise an alarm if the classification has a strong confidence in its identification.

Their evaluation considers either 10 or 6+2 ECUs and is tested in real car systems. The identification rate seems pretty good and they have no false-positives. However, I wonder how easy it might be to reduce the confidence in the classification; that appears like a potential vulnerability.

4.2 Detecting Attacks Against Robotic Vehicles: A Control Invariant Approach

Speaker: Hongjun Choi (Purdue University), collaboration with Wen-Chuan Lee (Purdue University), Yousra Aafer (Purdue University), Fan Fei (Purdue University), Zhan Tu (Purdue University), Xiangyu Zhang (Purdue University), Dongyan Xu (Purdue University), Xinyan Deng (Purdue University)

⁵I write that, because for our privacy buckets we do something similar and we order the privacy losses.

Robotic Vehicles (RV) include hobby drones, delivering drones, military UAC’s self driving cars. Abstractly, these RV’s are physical systems that interact with the physical world, while under control of a cyber-system. The main motivation for the talk is that while we are starting to understand and defend against cyber attacks, new attacks focus on physical attacks: new signals, spoofed sensors, or throwing stones at drones. A cute example is disturbing the gyro-sensors using noise (actual audio signals) that will influence the behavior of a drone.

To detect such attacks, they employ control invariants that, e.g., check whether the laws of physics are seemingly violated in the sensor data. To this end, they predict the next system state during runtime and compare this prediction with what they measure. They reverse-engineer the control graph of the RV and insert their invariants into the control program binary of the RV.

They show that their prediction is pretty close to the measurements; their errors are fairly small and at least all of their own attacks were detected. The authors are aware that this doesn’t yet make the system secure: They perform an attack and while they immediately detect it in their ground-based system, in the speaker’s words, “the drone still crashes, but we see an error message”.

4.3 Truth Will Out: Departure-Based Process-Level Detection of Stealthy Attacks on Control Systems

Speaker: Magnus Almgren (Chalmers University of Technology), collaboration with Wissam Aoudi (Chalmers University of Technology), Mikel Iturbe (Mondragon University)

The talk tackles industrial control systems and their vulnerabilities, e.g., the blackouts caused in Ukraine. Attacks on such systems can be devastating and we need to combine IT with operational technologies to defend against them. ICS systems often are cyclic and deterministic. Thus, “normal” behaviour can be learned or even modeled.

Ideas of previous work build a model of the physical process, then use the model to predict future system behavior and compare the predictions with the observations and raise alarms whenever the deviation is too large. Magnus declares that predicting the future is unnecessarily difficult and thus their PASAD system (Process Aware Stealth Attack Detection) focuses on solving an easier problem: They thus require only limited knowledge and also detect subtle and stealthy attacks. PASAD uses the raw features and is model-free.

PASAD works in two phases:

- Offline learning: They extract signals from the system, reduce noise, construct the signal subspace and project training vectors and compute a centroid to define the normal behavior.
- Online detection: They project the same vectors and try to see whether there is a deviation from the “normal” behavior. They raise an alarm whenever the deviation exceeds a threshold.

This is a nice safeguard, but again has the same limitations that the previous talks on this field had: we don’t get any kind of guarantee and it is not completely clear, how difficult it would be to fool the system. Still, I think it is a nice monitor that should make attacks more complicated.

4.4 On the Safety of IoT Device Physical Interaction Control

Speaker: Wenbo Ding (Clemson University), collaboration with Hongxin Hu (Clemson University)

Unsurprisingly, smart homes are still on the rise; the amount of devices grows fast and they obtain more and more complicated functions. The talk focuses on physical features and their impact on smart devices, e.g., smart windows might influence the smart heater control. The reverse direction is much more fun: If the attacker can hack into the heater, raise the temperature and thus force the temperature control application to open the window, thus allowing an attacker to open a window and thus break into a house, by hacking into the heater. As another example, an app might check the status of some sensor to lock the house if the owner leaves. Consequently, the app might lock or unlock the physical locks of the house, depending on sensor data.

To protect against such angles we first need to identify all physical angles and then their interactions. To this end, they first analyze their applications for intra-app issues and physical channels. From these two features they build an interaction chain and then they analyze the risk of these chains.

Analysis in more details In their intra-app analysis, they look for flows within the application. They also check which devices might be used and how they can effect each other. Next they try to identify all physical channels from the description using natural language processing. They parse the sentences into words, extract nouns and then infer channels. I’m not quite sure why they don’t analyze the sensors directly. They then generate interaction chains to figure out which snippets of interactions can be linked to form larger chains. Finally, they perform a risk analysis; as a base-line they use the intra-app interactions they have previously found (this is why they need them). They assume that intra-app interactions are safe, so they compare their physical interaction chains with the intra-app interaction chains to then classify the former as “probably safe” or “probably not safe”. In more details, they actually have a more fine-grained model of the normal behavior: they classify the intra-app chains into several types of chains, e.g., “temperature related stuff” and “movement related stuff”; if the physical interactions fall into the same classes, they are considered fine. If they fall outside of the classes learned on the intra-app interactions, the distance to the nearest class is the “risk factor” of these samples.

Looking at their evaluation, the method seems to leave a lot of room for improvement. It is better than pure random guessing, but not by too much.

5 CCS, Wednesday, Usable Security

5.1 Asking for a Friend: Evaluating Response Biases in Security User Studies

Speaker: Elissa M. Redmiles (University of Maryland), collaboration with Ziyun Zhu (University of Maryland), Sean Kross (University of California San Diego), Dhruv Kuchhal (Maharaja Agrasen Institute of Technology), Tudor Dumitras (University of Maryland), Michelle L. Mazurek (University of Maryland)

Motivation There is an increasing number of surveys published at the top security conferences. A key-question thus is whether the answers given by people questioned for these surveys are correct. There are many aspects, but Elissa focuses on the questions in this talk.

Evaluation / dataset For their evaluation, they use Symantec host records (500k people) on whether and how fast people updated their software and then performed a survey on 2k people how they would intend to respond to a message that there is a new update. They picked the answer choices to match the frequencies they observed in the data. They also phrased the question s.t. it matched previous work. They suggest that you always include an “I don’t know / don’t want to answer” in any survey, to prevent people from randomly answering if they don’t want to answer. They performed a significant number of pretesting steps, but Elissa emphasizes that pilots should only be used for making sure the technical stuff works.

Biases and lessons learned Not surprisingly, they found a systematic bias, i.e., they answer more positively when questioned whether they would update their system. When asking what they would recommend a friend should do instead of what they think they’d do, the answers are even more positive.

They tried to measure the “cost” of updating their system, including “having to restart the system” and “the observed number of crashes” (including observations on the first derivative of crashes: “does the application crash more or less often?”). They found that people who generally tended to update (as by their metric of costs) tended to answer that they intended to update more (significant positive effect); there also was a much smaller effect on whether they indeed updated faster.

To filter out subjects that answered wrong or illogical things, e.g., that claimed they wanted to update because the update didn’t require a restart if the message explicitly said a restart was required. Overall though, Elissa paints a bleak picture and suggests not using surveys for finding the actual truth on statistics, but more for getting an impression on “why” people act in certain ways.

5.2 Towards Usable Checksums: Automating the Integrity Verification of Web Downloads for the Masses

Speaker: Kévin Huguenin (UNIL – HEC Lausanne), collaboration with Mauro Cherubini (UNIL – HEC Lausanne), Alexandre Meylan (UNIL – HEC Lausanne), Bertil Chapuis (UNIL – HEC Lausanne), Mathias Humbert (Swiss Data Science Center, ETH Zurich and EPFL), Igor Bilogrevic (Google Inc.)

Motivation Kévin starts by cleverly confronting us with the fact that even we (as security researchers) don't actually check the checksums of software.

Checksums are often put on websites to allow users to verify that software has not been tampered with by an adversary. Since this has to be done manually, Kévin asks the rhetorical question of whether such a barbaric technique is still appropriate in 2018. He asks a few obvious research questions (do people use checksums and are there problems?), gives obvious answers and then improves the state-of-the-art by providing a novel tool for checking integrity.

Surveys:

- In a survey of 2,000 people we see that more than half the people do download software from vendor/developer websites (making them vulnerable to tampering, but also allowing for improvement); about a quarter of people remember to have seen checksums. They asked people for the purpose of checksums and find that about 5% of people know that.
- When checking which types of hashes are provided by websites, about half of the ones they looked at used insecure hashes (like MD5) and only a small number of websites included instructions for how to check that the checksum is correct. Finally, they
- They performed a small (n=40) study on how people actually act when confronted with checksums. Most would download software from websites, but only one third was aware of checksums. They then put the subjects in front of an eye-tracker while having them download a file, compute the checksum, check the checksum and then extract and note down some information about the program, e.g., the version number (a useless instruction used to make sure that participants were not too aware of the aim of the study). The checksum of the third program they downloaded was incorrect, but the beginning and end were correct (which I think is pretty mean): 38% of participants did not detect this mismatch, even though they were explicitly instructed to verify the checksum. This correlation rate was not correlated with prior knowledge about checksums. They also noticed that people mostly looked at the beginning of checksums, not at the end of checksums.

Solution Kévin admits that there already is something like a solution: with SRI (Subresource integrity) the creator of a website can include an integrity field into the `<script>` tag on the page, which will lead to the file not being downloaded if the checksum doesn't match. Consequently, they extended SRI to also work in an `<a>` tag to allow an easier inclusion into download links. Their browser extension extracts the checksums from `<a>` elements, computes the checksums of the downloaded file and displays a message indicating whether the checksums match. Moreover, they provide an extension for Wordpress.

5.3 Investigating System Operators' Perspective on Security Misconfigurations

Speaker: Tobias Fiebig (TU Delft), collaboration with Constanze Dietrich (Berliner Hochschule für Technik), Katharina Krombholz (CISPA Helmholtz Center (i.G.)), Kevin Borgolte (Princeton University)

Motivation Tobias mentions that misconfigurations are a major issue for security. In this work, they started with exploratory interviews using IRC (for some reason) and then performed a study using a questionnaire on about 200 system operators.

Selection of results Over 75% answered that they had made misconfigurations in general, but even 90% admitted that they had made a specific misconfiguration; with the exception of just one operator, everyone answered that they had encountered someone else making a misconfiguration.

They have encountered pretty terrible misconfigurations, including the combination of username `admin` with password `admin` and skipping updates. Things seem to go wrong because of a lack of knowledge, (allegedly not due to poor online resources), overwhelming responsibility (allegedly not due to insufficient funding), and using defaults (allegedly not due to unhelpful standards).

They asked operators on whether they think their managers know what they are doing. Non-IT and governmental OPs were more skeptical of their managers than OPs working in IT. As a funny side-note: trust in their own tools seems to be directly correlated with juniority of the operators, which I think makes sense. Moreover, although OPs like blameless post mortems, i.e., allowing people to honestly report their mistakes and not being punished as long as they were not completely careless, they also answered that their companies did not budget for errors.

5.4 Detecting User Experience Issues of the Tor Browser In The Wild

Speaker: Kevin Gallagher (New York University), collaboration with Sameer Patil (Indiana University Bloomington), Brendan Dolan-Gavitt (New York University), Damon McCoy (New York University), Nasir Memon (New York University)

Motivation We get a very brief background on Tor⁶, including the Tor browser, a by now well-known tool for browsing the web anonymously. The main research question is what the user’s naturalistic web browsing experience is for people that use the Tor browser.

The goals then are to study naturalistic Tor use, i.e., in a real-world setting and not in a lab, to also preserve the privacy of users, while finally get fine-granular data. I think they asked people to install the Tor browser on their own machines and then performed questionnaires, interviews and finally asked people to write about why they would bypass the Tor browser and use another browser instead.

They wrote a Python script to look for the launches and actions of the Tor browser. if the participant uses the Tor browser, they would change to the state “Tor browser open”; when the respective browser was closed, participants were asked to answer a questionnaire. If after opening a browser the participant opened another browser (without closing the first one), the state went to “browser switch” and they were asked about a respective “why did you switch” questionnaire after they closed the browser.

Dataset and results They only had a $n=19$ dataset, heavily biased towards young male participants; they collected about 120 questionnaires. A main issue they heard was that sites did not work properly, e.g., news sites did not properly work in the Tor browser. Moreover, participants were annoyed that they were treated differently (e.g., had to perform several captchas before being able to access search queries). Moreover, they complained that localization led to websites in foreign languages. In addition people reported a lack of convenience-features, such as lack of easy access to bookmarks and password managers.

On the positive side, participants enjoyed warnings, e.g., that websites could observe the size of the browser window (and the dangers of resizing) and the information on the Tor circuit through which their traffic was routed.

Lessons learned and conclusion Kevin observes that selecting an appropriate attacker model might be important or relevant for many users: not every user needs to be worried about every attack angle the Tor browser protects against.

Overall, while these results give us some additional insight into the user experiences of the Tor browser, a sample set of $n = 19$ means that we probably should not take these results as final, but rather as a bit of evidence; consequently, more studies, involving more participants and larger sample sizes are an interesting area for future work. Since their script-driven approach seems easily applicable, such studies might be possible without significant hurdles.

⁶<https://www.torproject.org>

6 CCS Thursday, Web Security 1

6.1 Predicting Impending Exposure to Malicious Content from User Behavior

Speaker: Mahmood Sharif (Carnegie Mellon University), collaboration with Jumpei Urakawa (KDDI Research), Nicolas Christin (Carnegie Mellon University), Ayumu Kubota (KDDI Research), Akira Yamada (KDDI Research)

They attempt to preemptively protect users by predicting exposure to malicious pages. Short-term within session prediction can include alerts and prioritizing traffic for expensive analyses, but also straight-out blocking of dangerous content or connections.

Data collection They collected http requests of 20k customers of KDDI, and asked them in an online survey to assess their security awareness and how they handled security incidents / their security knowledge / their security behavior. Finally, they used the Google Safe Browsing blacklist as a ground-truth to see whether users connected to unsafe pages.

The goal then is to observe behavior early in every session to predict whether they will be exposed later in the same session (continuous usage with less than 20 minute pauses in between). Mahmood notices that although only few sessions are dangerous (0.1%), about 11% of users are exposed at some point. Interestingly, they notice that the usage of dangerous webpages spikes just before these pages are listed on the GSB list. In general, exposed users engage in longer sessions, i.e., there is some correlation between longer usage and exposure. Exposed users request “certain topics”⁷ more often than unexposed users. Exposed users generally tended to be more active from afternoon to midnight and men are twice as likely to get exposed than women. Together with previous points there may be some interpretation here, but Mahmood doesn’t go there.⁸

Using the features, they train a neural network, e.g., if a user after visiting reddit.com visits streams.xyz and then becomes exposed, their network should learn that if streams are visited after reddit, the user is more likely to be exposed. They show with a nice graph that predicting the future for just a few seconds might be somewhat possible (75% true positive rate with 20% false positives, or just above 56% true positives with 3% false negatives); if they additionally use context features, their performance improves. There is some evidence that actually many more pages might be malicious than just the GSB list indicates, leading to about 2 true positives for every false positive.

6.2 Clock Around the Clock: Time-Based Device Fingerprinting

Speaker: Iskander Sanchez-Rola (Deustotech, University of Deusto), collaboration with Igor Santos (Deustotech, University of Deusto), Davide Balzarotti (Eurecom)

Motivation The talk is about device fingerprinting. The use might be malicious (an attacker wants to identify targets), a company that tries to identify machines for enforcement of legal rights. Advertisers moreover track users without storing information on the client side (to track their browsing history) and banks might use tracking to harden their verification process. The aim of the talk is to improve these fingerprinting techniques; this seems like a very attack-oriented talk and it is not clear how to protect against such fingerprinting.

Using clocks for fingerprinting The talk tries to improve existing fingerprinting techniques and found that the execution time of certain functions seems to be a good fingerprint. Existing work needs access to clock cycles, requires a sound card with its own internal crystal clock and requires an hour of runtime. They, however, manage to perform fingerprinting using the generic datetime api.

Results Their evaluation targeting JavaScript pseudo-random number generation includes computers with the same setup, components and software and they manage to distinguish with 100% accuracy any one of their 180 computers. In another test they used HTML5 CryptoFP stuffs for fingerprinting; they compared 256

⁷A strong example of this is seeing more ads, which makes sense since they can redirect users to more malicious pages

⁸Nikita Borisov, however, does not share his restraint and when asked about the impact of visiting adult content on exposure, Mahmood confirms the significance of this feature and that they indeed used it in their context.

identical machines (homogeneous setting) as well as 300 participants with their own machines (heterogeneous setting). In the homogeneous setting they only manage to identify 18% of computers (which is still better than the 0% of previous techniques). If they combine CryptoFP with canvas and WebGL, they managed to uniquely identify 80% of machines in the homogeneous group.

6.3 Web’s Sixth Sense: A Study of Scripts Accessing Smartphone Sensors

Speaker: **Gunes Acar (Princeton University)**, collaboration with Anupam Das (Carnegie Mellon University), Nikita Borisov (University of Illinois at Urbana-Champaign), Amogh Pradeep (Northeastern University)

Motivation There are lots of sensors available on smartphones, including gyroscopes, light sensors, accelerometers. They are available to any web page without permission. The speaker invites us to visit <https://sensor-js.xyz/demo> and promises not to store or use our sensor data.

The goal is to figure out which websites actually access these sites and what the risks exactly are. Existing works indicate that key logging, pin recovery, fingerprinting and geolocation is possible. Even recordings can be obtained using non-microphone sensors. In terms of fingerprinting, the slightest differences in physical differences lead to slightly different behavior, allowing phones to be identified.

Evaluation They checked Alexas top 100k sites and investigated them to find out what exactly they are doing. To this end they used the OpenWPM-mobile crawler that appears like a real mobile browser. For this work, they extended the WPM browser to a mobile browser setting. They measured the data from their real device and made sure their browser matches a real Firefox fingerprint. They also provide some sensor data, i.e., they triggered some sensor events with values extracted from a real phone / mobile browser. This study reveals that many sites (3.6k), including popular news sites, accessed sensors. The most aggressive culprits were add delivery companies. To check whether they obviously used the data, they used easily recognizable values and then checked the URL requests and payloads for these values, finding a few hits.

To figure out why these sensors are accessed, they tried to cluster using (400 binary features per script, e.g., whether certain JavaScript API calls were used). They then manually analyzed 3-5 scripts per cluster, which was more difficult since some of the script were obfuscated. According to this analysis, the most common usage non-surprisingly is tracking (fingerprinting, audience recognition), followed by fraud detection. A rare, but noteworthy mention is a crypto script using motion sensor data to improve their random number generator.

Countermeasures

- Ad blockers seemed not very useful (2%-9% blocking rate)
- Feature policy APIs of browsers that allow to specify which sensors can be accessed are helpful, but not yet widely available.
- W3C specifies that you should not allow access from insecure and cross-origin frames, but not all browsers seemed to follow this recommendation.
- User permissions could be queried or a visual indication that sensors are accessed for high-precision readings, but this might impact usability.
- Safe browsing modes might help here as well.

Overall it is not completely clear how the sensor data is used; fingerprinting might work better if the phones are not moving (e.g., placed on a table), but there certainly are angles.

7 CCS Thursday, Web Security 2

7.1 Mystique: Uncovering Information Leakage from Browser Extensions

Speaker: Quan Chen (North Carolina State University), collaboration with Alexandros Kapravelos (North Carolina State University)

Motivation We look at the privacy impact of browser extensions. As an example, the Web of Trust extension, used by more than a million users, requires users to allow access to all website contents. There was an article showing that and how the developers of this extension later sold data they had collected.

Extensions have access to the DOM tree and can inject JavaScript into visited pages. Using additional privileges, they can directly query for information such as the user's browsing history, have persistent storage or leak their information.

Method The authors crawl the Chrome web store for extensions, then analyze them. Their system uses a dynamic taint tracking analysis for JavaScript to then track the usage of tainted data. The main challenge is how to technically implement taint tracking in JavaScript. Since the V8 engine keeps the sourcecode of any scripts run in the pages, they can implement dynamic taint tracking by supplementing it with information they gain via static analysis.

Overall, Quan explains many generic aspects of information flow analyses:⁹ They build a data flow graph (DFG) by following the assignment expressions including objects of interest (i.e., if the right hand side of the assignments contain tainted values). To include control flow dependencies (implicit flows), they treat all tests within control flows (e.g., conditionals and loops) as the right hand side of an assignment and all internal assignments are included in the DFG. Moreover, they try to tackle leaks occurring implicitly in function calls.

Evaluation They analyzed 180k chrome extensions, resulting in 350 flagged extensions, 78% of which they confirmed as true positives. Among these true positives, the top 10 most popular extensions alone had more than 60 millions users. Interestingly, 9 out of 10 of these extensions were security related extensions, i.e., extensions that promised to improve users' security. Moreover, they found a library that leaked data and that was used in more than 300 extensions.

7.2 How You Get Bullets in Your Back: A Systematical Study about Cryptojacking in Real-world

Speaker: Geng Hong (Fudan University), collaboration with Zhemin Yang (Fudan University), Sen Yang (Fudan University), Lei Zhang (Fudan University), Yuhong Nan (Fudan University), Zhibo Zhang (Fudan University), Min Yang (Fudan University), Yuan Zhang (Fudan University), Zhiyun Qian (UC Riverside), Haixin Duan (Tsinghua University)

Motivation Cryptojacking is a rising threat; as an example, while browsing a seemingly simple website, a user's CPU might be busy mining cryptocurrencies. Initial cryptojacking attacks were brute-force and pretty obvious (e.g., they included keywords like `coinhive` in their code). Modern attacks limit CPU usage and try to remain stealthy to avoid detection. Moreover, they use code-obfuscation techniques to thwart keyword based analyses.

Method The nature of cryptocurrency mining, however, requires regular repeated hash-based computations to actually mine coins. Consequently Geng proposes hash-based profilers to detect cryptojacking webpages efficiently. Most normal websites spend less than 1% of their time hashing, whereas mining sites spend most of the time hashing and there are only a few known/useful hashing functions. Thus, they annotate the nine most popular hash libraries; in addition, there are regular repeated call stacks on these pages that they can take into account for detection.

⁹I'm not exactly sure how they leverage existing work on information flow and where their taint tracking exceeds previous work

Evaluation They focused on Alexas Top 100k websites (and subdomains) and followed some external links. They then extract features and try to classify the websites. Parts of the questions they ask is who is most responsible for cryptojacking, finding that half the samples are entertainment and adult websites, most of which provided pirated resources like free movies or cracked games.¹⁰

The impact they observed includes 10 million users per month with quite a lot of computational power (and thus potentially revenue). Most of the infrastructure is not perceptible to users, as they only see their own local impact (in terms of computational power and potentially battery draining). The most effective countermeasure so far (except for their own work) seems to be blacklists, but they only have a limited effect: most malicious samples disappear or update frequently (within only 9 days). Previous blacklists are only updated every 10-20 days, which might explain this behavior.

Geng rhetorically asks why cryptojacking is so popular and then mentions that, indeed, there are several companies that provide easy-to-use cryptojacking libraries, including CoinHive. Finally, it seems that most wallet IDs are only associated with a small number of samples, which I think does not mean that they actually belong to different entities – creating a large number of wallets and then laundering the coins seems most reasonable for such shady business.

As their approach only requires a rather short-time crawling study, Geng suggests that better countermeasures could be developed based on their technique. Moreover, this could be interesting for a plugin or otherwise client-side approach that might allow users to choose whether or not they want to give consent to cryptocurrency mining.

7.3 MineSweeper: An In-depth Look into Drive-by Cryptocurrency Mining and Its Defense

Speaker: Radhesh Krishnan Konoth (Vrije Universiteit Amsterdam), collaboration with Emanuele Vineti (Vrije Universiteit Amsterdam), Veelasha Moonsamy (Utrecht University), Martina Lindorfer (TU Wien), Christopher Kruegel (UC Santa Barbara), Herbert Bos (Vrije Universiteit Amsterdam), Giovanni Vigna (UC Santa Barbara)

Motivation The second talk about detecting cryptojacking again motivates the rise of cryptocurrencies that brings with it the threat of cryptojacking. Radhesh also mentions that and how existing defenses like CPU measurements and blacklisting don't work. Rhadesh shows how easy the usage of CoinHive is and how easily it can be configured to throttle the CPU usage.

In an interesting case, a media-player provided by one side included a miner and allowed its player to be embedded into other sites. I would consider that on the border of morality; they did not explicitly ask for consent (which is bad), but here I can see that allowing their player to be embedded in other sites can be a reason for asking for some return value.

Eval They analyze Alexas top 1 million websites, finding more than 1700 drive-by mining websites, crawl 3 internal pages each, visit each page for only 4 seconds (without giving consent to mining) and then analyze the several terabytes of raw data they collect from all these sites. First, they look for keywords associated with the CryptoNight algorithm. Moreover, they perform some filtering (Rhadesh refers to the paper), ending up with more than 1700 cryptojacking websites, about 700 of which were included on the main page. The speaker emphasizes evasion techniques like code obfuscation and in some cases even anti-debugging tricks, CPU-throttling.

To estimate the profits generated by cryptojacking, they use visitor statistics from SimilarWeb. They estimate how much computational power is provided by smartphones and laptops, eventually finding out that the most profitable sites make above 17k USD per month, with many sites making many hundreds or several thousand USD per month. Our media-player from before seems to total above 30k USD per month.

Countermeasures: MineSweeper The CryptoNight algorithm, used by the cryptojacking sites they analyzed uses five standard cryptographic functions, it is a memory hard algorithm. Their defense looks for

¹⁰Geng shows a list of mostly adult sites and suggests some of the audience might have contributed to the computational power.

the presence of these five crypto functions, then counts the number of loops in which these functions are called and then infers the algorithm run by the script. They use the `-dump-wasm-module` flag to instruct Chrome to log information about WASM (Web Assembly) usage, which they then analyze; in their evaluation where the large number of sites only used 40 unique scripts using these 5 functions, they flagged 36 of them as potentially malicious and confirmed that the remaining 4 were probably benign.

Additionally they propose that after analyzing WASM stuff, an additional analysis of CPU usage can provide some insights if one looks at what the CPU is actually used for most.

7.4 Pride and Prejudice in Progressive Web Apps: Abusing Native App-like Features in Web Applications

Speaker: Jiyeon Lee (KAIST), collaboration with Hayeon Kim (KAIST), Junghwan Park (KAIST), Insik Shin (KAIST), Soeul Son (KAIST)

Motivation Progressive Web Apps (PWA) make web apps more reliable, faster and more engaging; they can be used for offline browsing and they come with push notifications (previously only available to native apps). In their study they address the security and privacy risks to PWAs, including Cryptocurrency mining, inferring browsing history and using push notifications for phishing attacks.

A service worker is standard in HTML5, supported by major browsers and can even run when the browser is closed; Cache provides local storage; web push allows to re-engage users and can be received by service workers even when the browser is closed.

Evaluation They collected from Alexas top 100k, found that 4k used such features, mostly push. A website owner can customize the icon, the title and the body of the push notification, only the domain cannot be changed. Some browsers, such as Firefox and Samsung Internet do, however, not show the domain the push notification is coming from, which causes severe phishing risks.¹¹ Adding push notifications is easy and there are third-party libraries that make sending push requests easy and support features, including push requests for HTTP websites.

Jiyeon introduces a new push redirection attack: an adversary that can redirect can get users to fool users into allowing them push notifications. In The speaker emphasizes that in Web push the endpoint URL is confidential, since knowing this URL allows to send push notifications to the user. It seems that there are security options, but they are not widely used.

The cache aspect of PWA can be exploited to learn a user's browsing history. To this end, when the victim opens the attacking PWA offline an "on load" event will only be triggered if victims have visited target PWAs. This sounds like a lack of boundaries; still, Firefox 59.0.2 seems to be vulnerable to such a cross-domain attack.¹²

Finally, to stay within the tradition of this session of analyzing CoinHive, we see that service workers can be abused to continue mining even after the browser has been closed. To make this more viable (and to stop the service worker from falling asleep), they send push notifications to the service worker, which, again, seems possible even in complete secrecy, on the Firefox browser (it allows silent push notifications to sleeping service workers).

8 CCS, Thursday, Tor

8.1 Deep Fingerprinting: Undermining Website Fingerprinting Defenses with Deep Learning

Speaker: Payap Sirinam (Rochester Institute of Technology), collaboration with Mohsen Imani (University of Texas at Arlington), Marc Juarez (imec-COSIC KU Leuven, Belgium), Matthew Wright (Rochester Institute of Technology)

¹¹Chrome shows the domain.

¹²It is unclear to me why the PWAs would be allowed to even query for such information

Motivation We get a brief introduction into Tor,¹³ the most widely used anonymous communication protocol. In Tor, clients create a connection via three relays, called Tor nodes to communicate with any remote webserver. As is known, attackers that observe traffic between the client and the entry node, can perform website fingerprinting (WF) attacks.

Attackers can train a model by visiting a select set of websites. This small set of websites is called the monitored websites, whereas all non-targets are unmonitored websites. The attacker then tries to identify which websites are used when observing user traffic. Payap gives a brief overview over the pros and cons of closed world assumptions (good for attacks, not very realistic) and open world assumptions (harder to attack, probably more realistic)

To respond to effective website fingerprinting attacks is to add dummy packets, or to randomly delay packets. Two light-weight WF defenses pad large gaps without introducing any delays while only incurring a moderate bandwidth overhead ($\approx 50\%$). The Walkie-Talkie defense combines a bit of extra bandwidth and extra latency (about 30% each). Both these approaches reduce existing attacks quite a bit.

There already is some series of work on WF using machine learning, typically in closed world scenarios. In this work, the authors combine new ML methods with protection against defenses.

Deep fingerprinting Payap tries to convince us that deep learning can help learn deeper features and thus overcome existing defenses. Thus, the key difference to previous work AWF by Rimmer et al. is that this work uses a deeper model with more layers carefully designed for this specific task. The features are only the directions of traffic, i.e., ± 1 .

Evaluation Against non-defended datasets in closed world and open world scenarios, they achieve great accuracy (slightly above previous work). Against WTF-Pad-PAD where more than 60% of bandwidth is added, they still achieve about 90% accuracy and here start to significantly outperform previous work. Against Walkie-Talkie that perfectly mixes patterns of two websites together, they perform a top-2 prediction with above 98% accuracy.

Their approach can not necessarily be guaranteed, since the convergence of the model is not necessarily going to occur; however, since this is an attack paper, the existence of an attack should suffice. The speaker assures us that there was only a small variance in accuracy/precision/recall between several trained models.

8.2 Privacy-preserving Dynamic Learning of Tor Network Traffic

Speaker: Rob Jansen (U.S. Naval Research Laboratory), collaboration with Matthew Traudt (U.S. Naval Research Laboratory), Nicholas Hopper (University of Minnesota)

Motivation Tor allows you to protect your privacy and is an open research and development platform. There is a rich body of research and a major research area for many prominent universities and many MSC and doctoral theses focus on Tor.¹⁴

Tor research often depends on Tor experimentation tools, such as Rob’s Shadow network simulator and others. On the live Tor network there is real and realistic traffic, whereas simulations often only model sites as single files. Using just one single file fails to capture a wide variety of important measures. To do better, they use PrivCount [CCS’16] and then extend it to safely measure ground truth, then learn generative models of Tor traffic (packets and streams) using hidden Markov modeling and iterative measurements.

PrivCount has forward secrecy and only gives aggregate statistics that are made differentially private by adding noise.

Evaluation They deployed PrivCount on the public Tor network on a couple of dozen nodes and measured a number of statistics, e.g., the number of circuits created by how many clients, each for active (i.e., used) and inactive (i.e., predictively built) circuits.

Based on these measurements together with exit relay observations. Exit nodes can learn when circuits are opened and closed, when streams are started or ended. Their hidden Markov model tries to capture

¹³It’s supposed to be capitalized like this, not in all caps.

¹⁴Including Rob’s and my own

the behavior of active and inactive users, starting and ending circuits. They start with an openly available TCP dump of openly available data, then put it into their PrivCount process and then update the HMM’s probabilities using the weight parameters. They update their model over 14 epochs of one day each and find that their model initially improves (until day 9), but then deteriorates. I’m not quite sure why that would be the case. Moreover, they create several traffic generator models that, e.g., create TCP connections and transfer data.

To evaluate their model they updated Shadow using RIPE Atlas, making about 5 million pings, then use the PrivCount version of Tor, record PrivCount events, run event traces through local PrivCount deployment and compare to previously collected “ground truth”, also measured using PrivCount.¹⁵ Across all of the statistics they measured the PrivCount model is closer to the ground truth, when compared with the single file model and just the bare protocol model. However, these statistics also show that the traffic still looks quite different from the traffic observed in Tor.

8.3 DeepCorr: Strong Flow Correlation Attacks on Tor Using Deep Learning

Speaker: Amir Houmansadr (University of Massachusetts Amherst), collaboration with Milad Nasr (University of Massachusetts Amherst), Alireza Bahramali (University of Massachusetts Amherst)

Motivation After yet another brief introduction into Tor, Amir emphasizes the dangers of traffic correlation attacks. Potential adversaries that might run such attacks include compromised Tor relays, Internet Exchange Points (IXP), Autonomous Systems (AS), compromised routers, and DNS servers.¹⁶ Consequently, traffic flow correlation is key to many Tor attacks. Amir controversially claims that Tor does not care, because Global adversaries are excluded from Tor’s threat model and existing traffic correlation attacks supposedly provide poor performance on Tor at scale.¹⁷ Furthermore, Amir claims that existing techniques cannot capture Tor’s complex and non-linear behavior.

Methodology and evaluation In this work, they use deep learning to design a flow correlation system that has much better accuracy than previous work, as it can deal well with non-linearity. They learn a correlation function for Tor instead of using generic correlation functions.

They crawled Alexas top 50k websites, used half for training (5k for most experiments), created about 1k arbitrary Tor circuits and then waited for 6 weeks before evaluating their model. They used only 300 packets each, i.e., about 300KB of traffic and on those manage to achieve a significantly better accuracy when compared, e.g., with the RAPTOR attack. Their main aim is to learn a correlation function that is independent of the circuits it was trained on. They see that the model trained on different circuits, destinations and with a one-week gap, is able to perform almost as well as if it is trained on the same circuits and destinations and without a time-gap. According to Amir the system has to be retrained only about once per month to preserve the performance of the algorithm. As is to be expected, the more packets they use, the more stable are their results. Amir suggests the attacker could implement the attack in a hierarchical manner, allowing them to adaptively choose the model depending on the amount of observed packets. In terms of runtime, they require about one day of training time, but only 2ms for evaluating a sample.¹⁸

Amir ends the talk by mentioning countermeasures, including orthogonal countermeasures such as changing routing to be AS-aware¹⁹; pluggable transports might help counter their correlation attacks, but reduce the performance.

¹⁵I’m not completely sure here.

¹⁶All of these attacks are covered in my thesis, which btw. can be found here: https://publikationen.sulb.uni-saarland.de/bitstream/20.500.11880/26754/1/thesis_FINAL.pdf

¹⁷This statement seems to contradict several existing works, including the RAPTOR attack.

¹⁸It is quite possible that in order to really apply the technique one would have to look at a large number of pairs

¹⁹Changing routing depending on where the user is can expose users to other types of attacks; I refer to my thesis for an analysis.

8.4 Measuring Information Leakage in Website Fingerprinting Attacks and Defenses

Speaker: Nicholas Hopper (University of Minnesota), collaboration with Shuai Li (University of Minnesota), Huajun Guo (University of Minnesota)

Motivation Nick starts with talking about the first paper of the session: Website fingerprinting papers crawl Tor, then train classifiers on that data and then use them to classify stuff. There are several relevant questions, including the choice of features, of classifiers, of the “other” dataset, as well as how well the approach can scale to more pages/sites. Moreover, there seems to be no obvious way to incorporate auxiliary information and this might lead to typical classifier metrics understating the attack effectiveness.²⁰

Approach They want to model website fingerprints by their probability density functions. They propose the Website Fingerprint Density Estimation (WeFDE) framework for modeling FP density functions. They then use information leakage as a more complete evaluation of WF features and defenses.

They conducted a large-scale study, looking at 3k features used in the Tor website fingerprinting literature, based on a large dataset with 200k Tor web browsing visits and evaluated 6 WF defenses to show differences in accuracy and information leakage.

WeFDE collects the features, then computes the mutual information between pairs of features, clusters features and then works with them. A significant challenge here is to compare features with different densities, e.g., some might be point masses, whereas others are continuously distributed. Moreover, sometimes features can be highly correlated, while other times they are uncorrelated.

They fit Gaussians to their point features; if there are individual points, they fit really narrow Gaussians around them, whereas if there are more points, they combine them and fit wider Gaussians around them. I suspect they then create the PDF as a Gaussian mixture of these Gaussians. To combat the huge number of features they cluster them and reduce redundant features and as one might expect, there was a lot of redundancy in the features: even among the top 200 features, about half were completely redundant.

The information leakage measure suggests that some defenses might reduce the accuracy, but not necessarily the information leakage. It is not completely clear to me what exactly that means, so I would recommend having a look into their paper for an overview over their evaluation.

9 Post-conference Workshop AISEC

Remark: I am not a machine learning expert (yet); thus, these summaries might have an even lower accuracy and be shorter than previous ones. Feel free to contact me with any corrections.

9.1 Hate speech detection

Speaker: Tommi Gröndahl (Aalto University), collaboration with Luca Pajola (Aalto University), Mika Juuti (Aalto University), Mauro Conti (University of Padua) and N. Asokan (Aalto University) Hate speech is a real problem that has a significant negative social impact, but since it is a semantic property it is hard to automatically detect via machine learning.

Motivation Existing methods tend to treat hate speech as a standard text classification task, which The main aim of this work is to compare state-of-the-art classifiers on different datasets and to gain insights into challenges facing ML models and their resilience to a variety of attack types.

Method Datasets are Wikipedia edit comments and several sets of Twitter comments and types of hate speech range from sexism over racism to personal attacks.

They replicated 7 existing models and then compared their performances as two-class models. The models did comparably well when applied to the same testset as the dataset, but they performed badly when used

²⁰E.g., Walkie-Talkie can be defeated by looking at the top two results instead of the top result, as we’ve seen in the Talk by Payap.

on a testset from another set. Generally, the models seemed to be easily offended: they considered offensive statements as “hate speech”, even though they were not exactly considered hate speech by the researchers.

They then tried resilience to evasion attacks, such as typos (“I htae you”) or leet speech (“I h4t3 u”), adding whitespaces and otherwise perturbing the data or appending words (such as random english words or words specifically taken from the non-hate speech training set). Whitespace removal non-surprisingly defeats word based models, whereas character based models were not affected much. In word appending, the results were mixed.

The “love” attack Based on their results they performed the “Love” attack: they removed all whitespaces and then added the word “love”, e.g, turned “I hate you” into “Ihateyou love”. or “They are liberal idiots who are uneducated” into “TheyAreLiberalIdiotsWhoAreUneducated love”, thus reducing confidence in hate speech classification from 96% to 15%. Character models are significantly more resilient against such attacks, which makes sense since the whitespace removal defeats word-based models. Character-based models also performed comparably with deep neural networks, but again were resilient against such attacks.²¹

9.2 Towards Query Efficient Black-box Attacks: An Input-free Perspective

Authors: Yali Du (University of Technology, Sydney), Meng Fang (Tencent AI Lab), Jinfeng Yi (JD AI Research), Jun Cheng (Shenzhen Institutes of Advanced Technology, Chinese Academy of Sciences) and Dacheng Tao (The University of Sydney)

Motivation We first see how an adversarial example from a black-box attack might not be perceivable by a human, but can defeat Google’s Label Detection API quite severely. A similar attack makes Baidu’s Animal Detection API classify a gray square into an animal. Such attacks can be quite severe when face unlocking systems or fingerprint recognition systems can be fooled, thus allowing access to systems.

Input-free attacks turn gray images into adversarial examples and are particularly applicable when the adversary is free to choose any image as an input or has no knowledge of the real inputs. They are highly effective, but can easily be detected by humans.

Results They developed a novel such attack and compare it with other existing attacks. Their region-based attack is as successful as previous attacks, but requires significantly less queries and has a smaller L_∞ deviation.

9.3 Stochastic Substitute Training: A General Approach to Craft Adversarial Examples against Defenses which Obfuscate Gradients

Speaker: Greg Cusack (University of Colorado Boulder), collaboration with Mohammad Hashemi (University of Colorado Boulder) and Eric Keller (University of Colorado Boulder)

Motivation Following the motivation of both the invited talk and the previous talks in this session, we hear that neural networks are everywhere, but vulnerable to attacks. Greg adds to this that even modern robust neural networks, designed specifically to defeat such attacks, are vulnerable to white-box attacks. The goal of this talk now is to extend these attacks to purely black-box attacks that require only minimal access

Existing black-box attacks are very effective against vanilla models, but seem not to be robust in the presence of modern defenses: Fortifying defenses try to predict the correct class of adversarial examples (attackers can send inputs and see logits); detecting defenses try to identify the presence of adversarial examples.

²¹A slight critique from the audience should be mentioned: character-based attacks might defeat character-based attacks, but affect word-based less.

Method They first augment their dataset with random noise, then train a substitute model, querying the robust model for labels. They then craft adversarial examples against their substitute model and use an optimizer to minimize their loss function. As soon as they find an adversarial example, they iterate to find an example that has a smaller perturbation. They consider a range of specific defenses and achieve a great success rate (100% according to Greg).

- Against the Thermometer defense, they create several substitute models at once and use all of them together to create adversarial examples, thus managing to create adversarial examples with an even smaller perturbation (in comparison with using only one substitute model).
- Against the SafetyNet defense, they also had to train a substitute model for the detector, starting with minimal assumptions on where and how the detector is applied. Metrics for success here are to fool the classifier, while also having a significant confidence and fooling the detector as well.
- Finally, they defeat the Defense-GAN, which they managed on the second attempt.

Overall, they can craft adversarial examples without knowledge of the particular type of defenses.²²

9.4 Adaptive Grey-Box Fuzz-Testing with Thompson Sampling

Speaker: Siddharth Karamcheti (Bloomberg), collaboration with Gideon Mann (Bloomberg) and David Rosenberg (Bloomberg)

Motivation We get some background on fuzzing, i.e., automated software testing using randomized, ideally unexpected inputs; they can be black-box or white-box; gray-box fuzzers are mutating existing inputs to check code paths and they are fast and efficient in that they find relevant bugs. While existing algorithms such as AFL (American Fuzzy Loop) are good, they are inefficient and fairly random. Thus, Siddharth asks how we can improve these techniques with data-driven techniques. As an example: which type of mutation is successful for a given program might depend a lot on the program.

Method The main idea is to learn which are the right parameters for a given program. The probabilities for the mutators start out as unobserved, but then are optimized: they pick a mutator, collect data for a fixed interval, then count how often the mutator generated a “successful” input.²³

they tested their approach on a set of 200 binaries, each of which has a real “bug” added by a human user. Moreover, they analyzed a set of 4 binaries from Coreutils, each of which is injected with 100 synthetic bugs. They compare three different baselines: AFL (Vanilla), AFL in Havoc mode, and their own adaptive machine learning approach. The latter beats all former approaches by a significant margin, indicating that learning which mutator is useful helps. Moreover, this approach also found almost twice as many unique inputs that result in a crash. Interestingly, there are cases in which the vanilla AFL approach performed better, which means that their adaptive AFL method might have a niche, but is not always superior.

Overall, their method introduces a simple form of adaptive, learned fuzzing, which might help to find bugs in some cases. I’m not quite sure what it means to learn which mutators perform best; the author mentions that this might lead their approach to get stuck in a particular type (or a few types) of mutators. Still, their approach seems to perform better than the static/uniformly random ones, so this might be an interesting direction.

9.5 Toward Smarter Vulnerability Discovery Using Machine Learning

Speaker: Gustavo Grieco (Trail of Bits), collaboration with Artem Dinaburg (Trail of Bits)

²²Since this work focused mostly on MNIST and used CIFAR-10 for a few evaluations, an interesting direction for future work is to look at how these techniques perform on more complex datasets.

²³Success here means that the mutator discovered a new code path.

Motivation Similar to the previous talk, their aim as well is to explore how to best find inputs that explore program vulnerabilities. Similar to the previous talk, we first need to find good inputs, but then need to find good parameters for the fuzzer(s) we want to use.

Gustavo briefly repeats which types of fuzzers exist; the aim now is to use these tools to analyze a binary. Instead of proposing a novel technique in and of itself, their goal, however, is to find the best tool for the job for each target program and initial input (that can then, e.g., be mutated). We thus try to optimize the number of branches and ideas explored by the system before it manages to invoke a crash.

Data collection and training Their contribution is to collect and label a new dataset (288 binaries of different complexities), to then extract features from data they collected during execution (using Manticore to get, e.g., the number of system calls) that they then use to train several ML classifiers to predict which choices give better results and finally they develop another level of abstraction that supposedly can learn how to best approach this problem.

Their model chooses one of three fuzzers (Manticore, Grr and AFL) together with the suggested parameters for that fuzzer. For training they consider three labels: failure (no new program state explored or something didn't work), new input (new state explored), and valuable input (input led to a crash).

Their results indicate that they can indeed, with an accuracy better than pure guessing,²⁴ predict whether or not an action will give good results.

²⁴The accuracy was about 60%-70%, where pure guessing would give us 50% accuracy if the decision is unbiased; I'm actually not sure whether this assumption is valid.